

E1

T1.1

ProductName: macOS

ProductVersion: 15.0.1

BuildVersion: 24A348

Darwin Mac 24.0.0 Darwin Kernel Version 24.0.0: Tue Sep 24 23:37:13 PDT 2024;

root:xnu-11215.1.12~1/RELEASE_ARM64_T8112 arm64

time.time: 7.152557373046875e-07

timeit.default_timer: 8.297502063214779e-08

time.time_ns: 768.0

ProductName: macOS

ProductVersion: 14.6.1

BuildVersion: 23G93

Darwin Mac 23.6.0 Darwin Kernel Version 23.6.0: Mon Jul 29 21:14:30 PDT 2024;

root:xnu-10063.141.2~1/RELEASE_ARM64_T6030 arm64

time.time: 7.152557373046875e-07

timeit.default_timer: 8.288770914077759e-08

time.time_ns: 768.0

clock_granularity.py was used to test this which is located in

A1/E1/T1_1/clock_granularity.py

T1.2

We ran the command:

```
"python3 -m timeit -n 5 -r 1 -s "import JuliaSet"
```

```
"JuliaSet.calc_pure_python(desired_width=1000,
```

```
max_iterations=300) "
```

With the JuliaSet.py using our own made decorator.py in A1/E1/T1_2/decorator.py

Results for computer with core speed of 4.05 GHz (assumed):

Function: calculate_z_serial_purepython

Average Execution Time: 2.619 seconds

Average Standard Deviation: 0.0103 seconds

Function: calc_pure_python

Average Execution Time: 2.811 seconds

Average Standard Deviation: 0.0215 seconds

With a clock frequency of about 4.05 GHz, a single cycle is about 0.25ns which is vastly smaller than both standard deviations (as they have 10.3 ms and 21.5 ms respectively). This means the amount of cycles must differ between each run or other programmes take up cycles on the system when running the code.

It might as well be the case that the OS decides to not use the performance cores in the CPU. These have a lower clock frequency of 2.75 GHz which results in the per cycle time of

0.36ns. This is still magnitudes smaller than the standard deviation of 10.3 ms and 21.5 ms of the average standard deviations. The kernel thread scheduling could allocate varying amounts of CPU time to each thread, contributing to the observed variation in execution times.

T1.3

Using the cprofile command:

```
python3 -m cProfile -s cumulative JuliaSet.py
```

The top results with cProfile:

Ordered by: cumulative time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
118/1	0.003	0.000	5.366	5.366	{built-in method builtins.exec}
1	0.012	0.012	5.366	5.366	JuliaSet.py:1(<module>)
1	0.327	0.327	5.189	5.189	JuliaSet.py:27(calc_pure_python)
1	3.325	3.325	4.729	4.729	JuliaSet.py:67(calculate_z_serial_purepython)
34219980	1.404	0.000	1.404	0.000	{built-in method builtins.abs}
14	0.001	0.000	0.403	0.029	__init__.py:1(<module>)
157/1	0.001	0.000	0.165	0.165	<frozen importlib._bootstrap>:1349(_find_and_load)
157/1	0.000	0.000	0.165	0.165	<frozen importlib._bootstrap>:1304(_find_and_load_unlocked)
382/3	0.000	0.000	0.165	0.055	<frozen importlib._bootstrap>:480(_call_with_frames_removed)
145/3	0.000	0.000	0.165	0.055	<frozen importlib._bootstrap>:911(_load_unlocked)
116/1	0.000	0.000	0.165	0.165	<frozen importlib._bootstrap_external>:989(exec_module)
1	0.000	0.000	0.165	0.165	decorator.py:1(<module>)
375/10	0.000	0.000	0.159	0.016	{built-in method builtins.__import__}
251/41	0.000	0.000	0.158	0.004	<frozen importlib._bootstrap>:1390(_handle_fromlist)
2005342	0.132	0.000	0.132	0.000	{method 'append' of 'list' objects}
1	0.000	0.000	0.086	0.086	__config__.py:1(<module>)
116	0.001	0.000	0.058	0.000	<frozen importlib._bootstrap_external>:1062(get_code)
145/141	0.000	0.000	0.037	0.000	<frozen importlib._bootstrap>:806(module_from_spec)
25/24	0.000	0.000	0.035	0.001	<frozen importlib._bootstrap_external>:1287(create_module)
25/24	0.021	0.001	0.035	0.001	{built-in method _imp.create_dynamic}
156	0.001	0.000	0.033	0.000	<frozen importlib._bootstrap>:1240(_find_spec)
1	0.000	0.000	0.029	0.029	numerictypes.py:1(<module>)
116	0.000	0.000	0.028	0.000	<frozen importlib._bootstrap_external>:1183(get_data)
1	0.000	0.000	0.027	0.027	index_tricks.py:1(<module>)
116	0.024	0.000	0.024	0.000	{method 'read' of '_io.BufferedReader' objects}
1	0.000	0.000	0.024	0.024	defmatrix.py:1(<module>)
1	0.000	0.000	0.022	0.022	type_aliases.py:1(<module>)

cProfile is showing results higher than the built decorator because of the fact that it has a lot more overhead done than the decorator. It also differentiates calc_pure_python from calculate_z_serial_purepython part which means it does not include it in the first function in cProfile while the decorator includes both. This makes cProfile include more overhead but is as well more informative than the decorator.

Using the command

```
python3 -m kernprof -l JuliaSet.py
```

The results with line profiler becomes:

```
Timer unit: 1e-06 s
Total time: 6.04551 s
File: JuliaSet.py
Function: calc_pure_python at line 27

Line #      Hits          Time    Per Hit   % Time  Line Contents
=====
  27          1           2.0        2.0     0.0      @profile
  28          1           1.0        1.0     0.0      def calc_pure_python(desired_width, max_iterations):
  29          1           1.0        1.0     0.0          """Create a list of complex coordinates (zs) and complex parameters (cs),
  30          1           1.0        1.0     0.0          build Julia set"""
  31          1           1.0        1.0     0.0          x_step = (x2 - x1) / desired_width
  32          1           1.0        1.0     0.0          y_step = (y1 - y2) / desired_width
  33          1           1.0        1.0     0.0          x = []
  34          1           5.0        5.0     0.0          y = []
  35          1           0.0        0.0     0.0          ycoord = y2
  36       1001       1118.0        1.1     0.0          while ycoord > y1:
  37       1000       1602.0        1.6     0.0              y.append(ycoord)
  38       1000       1396.0        1.4     0.0              ycoord -= y_step
  39          1           4.0        4.0     0.0          xcoord = x1
  40       1001       1565.0        1.6     0.0          while xcoord < x2:
  41       1000       1564.0        1.6     0.0              x.append(xcoord)
  42       1000       1477.0        1.5     0.0              xcoord += x_step
  43          1           0.0        0.0     0.0          # build a list of coordinates and the initial condition for each cell.
  44          1           0.0        0.0     0.0          # Note that our initial condition is a constant and could easily be removed,
  45          1           0.0        0.0     0.0          # we use it to simulate a real-world scenario with several inputs to our
  46          1           0.0        0.0     0.0          # function
  47          1           0.0        0.0     0.0          zs = []
  48          1           0.0        0.0     0.0          cs = []
  49       1001       187.0        0.2     0.0          for ycoord in y:
  50       1001000    164862.0        0.2     2.7              for xcoord in x:
  51       1000000    160594.0        0.2     2.7                  zs.append(complex(xcoord, ycoord))
  52       1000000    158345.0        0.2     2.6                  cs.append(complex(c_real, c_imag))
  53          1           0.0        0.0     0.0          # print("Length of x:", len(x))
  54          1           0.0        0.0     0.0          # print("Total elements:", len(zs))
  55          1           0.0        0.0     0.0          start_time = time.time()
  56          1       5552752.0    6e+06    91.8          output = calculate_z_serial_purepython(max_iterations, zs, cs)
  57          1           2.0        2.0     0.0          end_time = time.time()
  58          1           2.0        2.0     0.0          secs = end_time - start_time
  59          1           2.0        2.0     0.0          # print(calculate_z_serial_purepython.__name__ + " took", secs, "seconds")
  60          1           0.0        0.0     0.0          # This sum is expected for a 1000^2 grid with 300 iterations
  61          1           0.0        0.0     0.0          # It ensures that our code evolves exactly as we'd intended
  62          1           0.0        0.0     0.0          # assert sum(output) == 33219980
  63          1           0.0        0.0     0.0          return output
  64          1           0.0        0.0     0.0          """

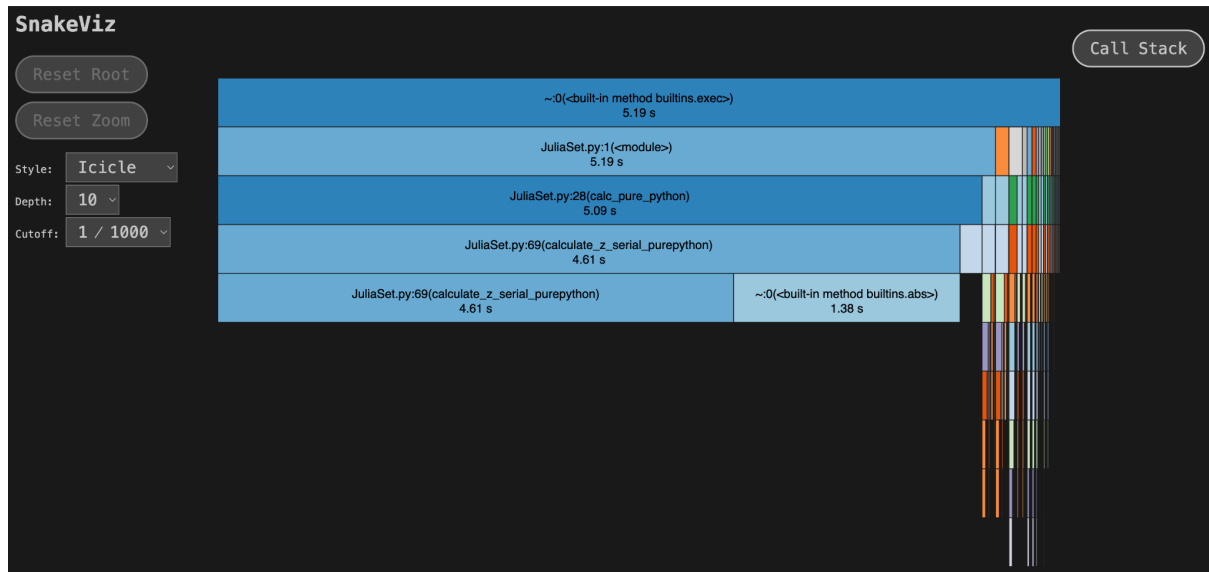
6.05 seconds - JuliaSet.py:27 - calc_pure_python
```

```
marcuscarlbom@MBPsomtrMarcus2 A1 % python3 -m line_profiler -rmt "JuliaSet.py.lprof"
Timer unit: 1e-06 s
Total time: 12.4048 s
File: JuliaSet.py
Function: calculate_z_serial_purepython at line 68

Line #      Hits          Time    Per Hit   % Time  Line Contents
=====
  68          1           0.0        0.0     0.0      @profile
  69          1           0.0        0.0     0.0      def calculate_z_serial_purepython(maxiter, zs, cs):
  70          1           0.0        0.0     0.0          """Calculate output list using Julia update rule"""
  71          1       965.0        965.0     0.0          output = [0] * len(zs)
  72     1000001    131045.0        0.1     1.1          for i in range(len(zs)):
  73     1000000    94417.0        0.1     0.8              n = 0
  74     1000000    94801.0        0.1     0.8              z = zs[i]
  75     1000000    87967.0        0.1     0.7              c = cs[i]
  76     34219980    5457886.0        0.2    44.0              while abs(z) < 2 and n < maxiter:
  77     33219980    3467934.0        0.1    28.0                  z = z * z + c
  78     33219980    2972864.0        0.1    24.0                  n += 1
  79     1000000    96916.0        0.1     0.8              output[i] = n
  80          1           1.0        1.0     0.0          return output

12.40 seconds - JuliaSet.py:68 - calculate_z_serial_purepython
```

Snakeviz visualisation



Measurements of overhead:

Running JuliaSet.py:

```
● marcuscarlbom@MBPsomtrMarcus2 A1 % python3 JuliaSet.py
This run took a total of 2.316828966140747 seconds
● marcuscarlbom@MBPsomtrMarcus2 A1 % python3 JuliaSet.py
This run took a total of 2.3311450481414795 seconds
● marcuscarlbom@MBPsomtrMarcus2 A1 % python3 JuliaSet.py
This run took a total of 2.3112940788269043 seconds
● marcuscarlbom@MBPsomtrMarcus2 A1 % python3 JuliaSet.py
This run took a total of 2.308480978012085 seconds
● marcuscarlbom@MBPsomtrMarcus2 A1 % python3 JuliaSet.py
This run took a total of 2.3560657501220703 seconds
○ marcuscarlbom@MBPsomtrMarcus2 A1 %
```

Average time 2.325 seconds

Running JuliaSet.py with cProfiler:

```
● marcuscarlbom@MBPsomtrMarcus2 A1 % python3 -m cProfile -o profile.stats JuliaSet.py
This run took a total of 5.087873935699463 seconds
● marcuscarlbom@MBPsomtrMarcus2 A1 % python3 -m cProfile -o profile.stats JuliaSet.py
This run took a total of 5.055094003677368 seconds
● marcuscarlbom@MBPsomtrMarcus2 A1 % python3 -m cProfile -o profile.stats JuliaSet.py
This run took a total of 5.107017993927002 seconds
● marcuscarlbom@MBPsomtrMarcus2 A1 % python3 -m cProfile -o profile.stats JuliaSet.py
This run took a total of 5.09935188293457 seconds
● marcuscarlbom@MBPsomtrMarcus2 A1 % python3 -m cProfile -o profile.stats JuliaSet.py
This run took a total of 5.090599298477173 seconds
○ marcuscarlbom@MBPsomtrMarcus2 A1 %
```

Average time 5.088 seconds, about 2.763 seconds of extra overhead which is about 219% increase in run time.

Running JuliaSet.py with line_profiler:

```

● marcuscarlbom@MBPsomtrMarcus2 A1 % python3 -m kernprof -l JuliaSet.py

This run took a total of 26.036913 seconds
Wrote profile results to JuliaSet.py.lprof
Inspect results with:
python3 -m line_profiler -rmt "JuliaSet.py.lprof"
● marcuscarlbom@MBPsomtrMarcus2 A1 % python3 -m kernprof -l JuliaSet.py

This run took a total of 26.394208 seconds
Wrote profile results to JuliaSet.py.lprof
Inspect results with:
python3 -m line_profiler -rmt "JuliaSet.py.lprof"
● marcuscarlbom@MBPsomtrMarcus2 A1 % python3 -m kernprof -l JuliaSet.py

This run took a total of 27.125499 seconds
Wrote profile results to JuliaSet.py.lprof
Inspect results with:
python3 -m line_profiler -rmt "JuliaSet.py.lprof"
● marcuscarlbom@MBPsomtrMarcus2 A1 % python3 -m kernprof -l JuliaSet.py

This run took a total of 25.749900 seconds
Wrote profile results to JuliaSet.py.lprof
Inspect results with:
python3 -m line_profiler -rmt "JuliaSet.py.lprof"
● marcuscarlbom@MBPsomtrMarcus2 A1 % python3 -m kernprof -l JuliaSet.py

This run took a total of 25.469858 seconds
Wrote profile results to JuliaSet.py.lprof
Inspect results with:
python3 -m line_profiler -rmt "JuliaSet.py.lprof"
○ marcuscarlbom@MBPsomtrMarcus2 A1 % █

```

Average time 26.15 seconds, about 23.825 seconds of extra overhead which is about 11247% increase in run time.

T1.4

Results from running

python3 -m memory_profiler JuliaSet.py

marcuscarlbom@MBPsomtrMarcus2 A1 % python3 -m memory_profiler JuliaSet.py
This run took a total of 1578.771924 seconds
Filename: JuliaSet.py

Line #	Mem usage	Increment	Occurrences	Line Contents
28	64.375 MiB	64.375 MiB	1	@profile
29				def calc_pure_python(desired_width, max_iterations):
30				"""Create a list of complex coordinates (zs) and complex parameters (cs),
31				build Julia set"""
32	64.375 MiB	0.000 MiB	1	x_step = (x2 - x1) / desired_width
33	64.375 MiB	0.000 MiB	1	y_step = (y1 - y2) / desired_width
34	64.375 MiB	0.000 MiB	1	x = []
35	64.375 MiB	0.000 MiB	1	y = []
36	64.375 MiB	0.000 MiB	1	ycoord = y2
37	64.406 MiB	0.031 MiB	1001	while ycoord > y1:
38	64.406 MiB	0.000 MiB	1000	y.append(ycoord)
39	64.406 MiB	0.000 MiB	1000	ycoord += y_step
40	64.406 MiB	0.000 MiB	1	xcoord = x1
41	64.438 MiB	0.031 MiB	1001	while xcoord < x2:
42	64.438 MiB	0.000 MiB	1000	x.append(xcoord)
43	64.438 MiB	0.000 MiB	1000	xcoord += x_step
44				# build a list of coordinates and the initial condition for each cell.
45				# Note that our initial condition is a constant and could easily be removed,
46				# we use it to simulate a real-world scenario with several inputs to our
47				# function
48	64.438 MiB	0.000 MiB	1	zs = []
49	64.438 MiB	0.000 MiB	1	cs = []
50	140.344 MiB	0.000 MiB	1001	for ycoord in y:
51	140.344 MiB	0.000 MiB	1001000	for xcoord in x:
52	140.344 MiB	22.828 MiB	1000000	zs.append(complex(xcoord, ycoord))
53	140.344 MiB	53.078 MiB	1000000	cs.append(complex(c_real, c_imag))
54				
55				# print("Length of x:", len(x))
56				# print("Total elements:", len(zs))
57	140.344 MiB	0.000 MiB	1	start_time = time.time()
58	165.375 MiB	165.375 MiB	1	output = calculate_z_serial_purepython(max_iterations, zs, cs)
59	165.375 MiB	0.000 MiB	1	end_time = time.time()
60	165.375 MiB	0.000 MiB	1	secs = end_time - start_time
61				# print(calculate_z_serial_purepython.__name__ + " took", secs, "seconds")
62				
63				# This sum is expected for a 1000^2 grid with 300 iterations
64				# It ensures that our code evolves exactly as we'd intended
65				# assert sum(output) == 33219980

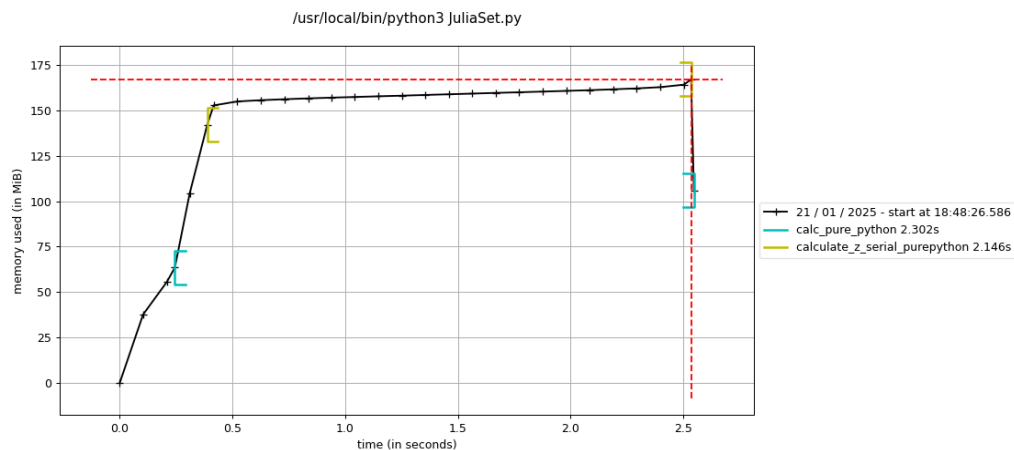
Filename: JuliaSet.py

Filename: JuliaSet.py

Line #	Mem usage	Increment	Occurrences	Line Contents
69	140.344 MiB	140.344 MiB	1	@profile
70				def calculate_z_serial_purepython(maxiter, zs, cs):
71				"""Calculate output list using Julia update rule"""
72	147.984 MiB	7.641 MiB	1	output = [0] * len(zs)
73	165.375 MiB	0.000 MiB	1000001	for i in range(len(zs)):
74	165.375 MiB	0.000 MiB	1000000	n = 0
75	165.375 MiB	7.156 MiB	1000000	z = zs[i]
76	165.375 MiB	7.156 MiB	1000000	c = cs[i]
77	165.375 MiB	0.000 MiB	34219980	while abs(z) < 2 and n < maxiter:
78	165.375 MiB	0.000 MiB	33219980	z = z * z + c
79	165.375 MiB	3.078 MiB	33219980	n += 1
80	165.375 MiB	0.000 MiB	1000000	output[i] = n
81	165.375 MiB	0.000 MiB	1	return output

marcuscarlbom@MBPsomtrMarcus2 A1 %

mprof plotting



Running the memory profiler 1 time took 1578.77 seconds. This is an increase of 1576.445 seconds which is 679 times longer than just running it.

Running the mprof command

```
python3 -m mprof run JuliaSet.py
```

Took on average after 5 runs 2.337 seconds. This is an increase of 0.002 seconds which is about an 0.5% increase in run time. Which due to there being deviations means it is likely having close to negligible impact on run time.

E2

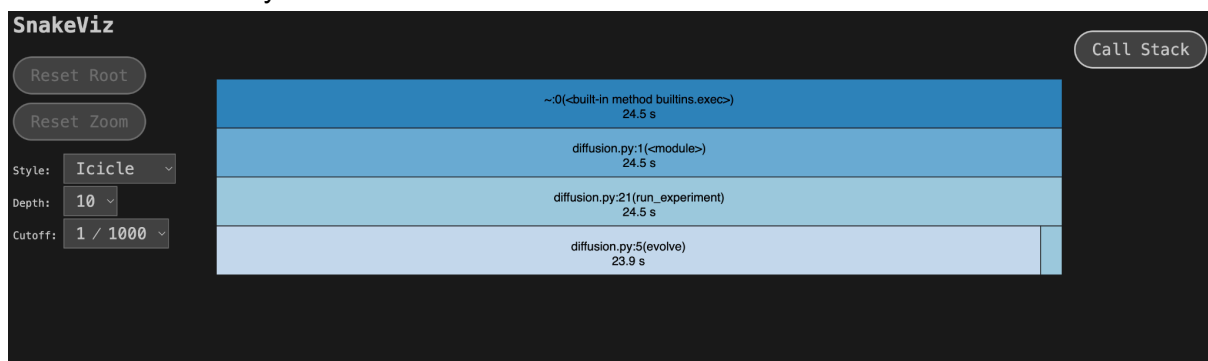
T2.1

```
marcuscarlbom@MBPsomtrMarcus2 E2 % python3 -m cProfile -s cumulative diffusion.py
304 function calls in 24.634 seconds

Ordered by: cumulative time

ncalls  tottime  percall  cumtime  percall filename:lineno(function)
1      0.000    0.000   24.634   24.634 {built-in method builtins.exec}
1      0.003    0.003   24.634   24.634 diffusion.py:1(<module>)
1      0.631    0.631   24.631   24.631 diffusion.py:20(run_experiment)
300    24.000    0.080   24.000    0.080 diffusion.py:4(evolve)
1      0.000    0.000    0.000    0.000 {method 'disable' of '_lsprof.Profiler' objects}
```

cProfile visualised by SnakeViz



Result of line_profiler:

```

marcuscarlbom@MBPsomtrMarcus2 E2 % python3 -m line_profiler diffusion.py.lprof
Timer unit: 1e-06 s

Total time: 89.5921 s
File: diffusion.py
Function: evolve at line 4

Line #      Hits          Time Per Hit   % Time  Line Contents
=====
4          300          184.0     0.6     0.0      @profile
5          192300        103921.0     0.5     0.1      def evolve(grid, dt, D=1.0):
6          192300        26586.0     0.1     0.0          xmax, ymax = grid_shape
7          123072000    14349472.0     0.1    16.0          new_grid = [[0.0] * ymax for x in range(xmax)]
8          122880000    10791285.0     0.1    12.0          for i in range(xmax):
9          122880000    20288001.0     0.2    22.6              for j in range(ymax):
10              grid_xx = (
11                  grid[(i + 1) % xmax][j] + grid[(i - 1) % xmax][j] - 2.0 * grid[i][j]
12              )
13              grid_yy = (
14                  grid[i][(j + 1) % ymax] + grid[i][(j - 1) % ymax] - 2.0 * grid[i][j]
15              )
16              new_grid[i][j] = grid[i][j] + D * (grid_xx + grid_yy) * dt
17          300          369.0     1.2     0.0          return new_grid

Total time: 178.91 s
File: diffusion.py
Function: run_experiment at line 20

Line #      Hits          Time Per Hit   % Time  Line Contents
=====
20          1          0.0     0.0     0.0      @profile
21          641          812.0     1.3     0.0      def run_experiment(num_iterations):
22          1          0.0     0.0     0.0          # Setting up initial conditions
23          641          812.0     1.3     0.0          xmax, ymax = grid_shape
24          1          0.0     0.0     0.0          grid = [[0.0] * ymax for x in range(xmax)]
25
26          # These initial conditions are simulating a drop of dye in the middle of our
27          # simulated region
28          block_low = int(grid_shape[0] * 0.4)
29          block_high = int(grid_shape[0] * 0.5)
30          for i in range(block_low, block_high):
31              for j in range(block_low, block_high):
32                  grid[i][j] = 0.005
33
34          # Evolve the initial conditions
35          for i in range(num_iterations):
36              grid = evolve(grid, 0.1)

```

Memory_profiler run on diffusion.py

```

marcuscarlbom@Mac E2 % python3 -m memory_profiler diffusion.py
Filename: diffusion.py

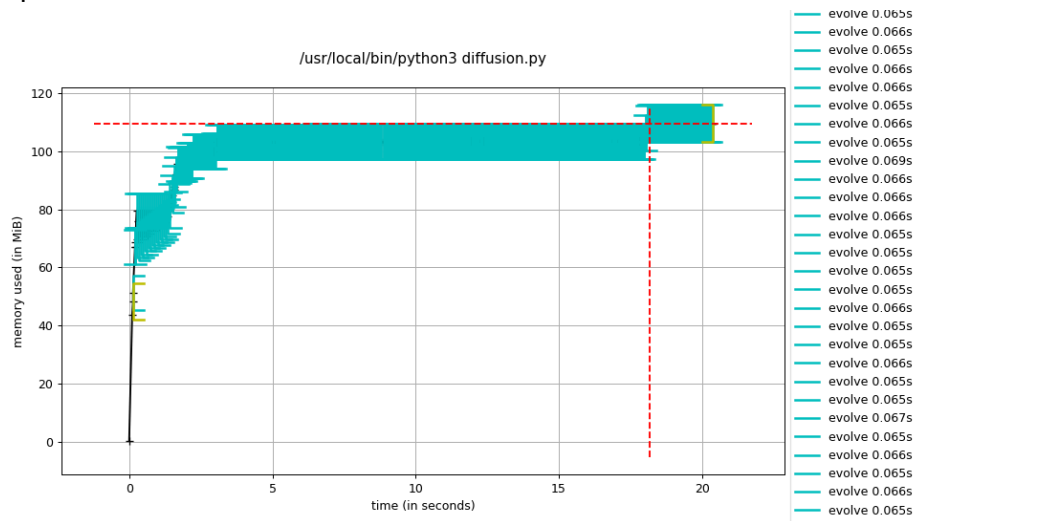
Line #      Mem usage      Increment  Occurrences  Line Contents
=====
4      93.812 MiB      21228.875 MiB          300      @profile
5      93.812 MiB      -6734.766 MiB          300      def evolve(grid, dt, D=1.0):
6      93.812 MiB      -4309864.156 MiB        192300          xmax, ymax = grid_shape
7      97.812 MiB      -5080139.344 MiB        192300          new_grid = [[0.0] * ymax for x in range(xmax)]
8      97.812 MiB      -3252085584.266 MiB      123072000          for i in range(xmax):
9      97.812 MiB      -3247012029.906 MiB      122880000              for j in range(ymax):
10              grid_xx = (
11                  grid[(i + 1) % xmax][j] + grid[(i - 1) % xmax][j] - 2.0 * grid[i][j]
12              )
13              grid_yy = (
14                  grid[i][(j + 1) % ymax] + grid[i][(j - 1) % ymax] - 2.0 * grid[i][j]
15              )
16              new_grid[i][j] = grid[i][j] + D * (grid_xx + grid_yy) * dt
17          300          7887.250 MiB          300          return new_grid

Filename: diffusion.py

Line #      Mem usage      Increment  Occurrences  Line Contents
=====
20      50.984 MiB       50.984 MiB           1      @profile
21      50.984 MiB       0.000 MiB           1      def run_experiment(num_iterations):
22      53.953 MiB       2.969 MiB          641          # Setting up initial conditions
23          53.953 MiB       0.000 MiB           1          xmax, ymax = grid_shape
24          53.953 MiB       0.000 MiB          641          grid = [[0.0] * ymax for x in range(xmax)]
25
26          # These initial conditions are simulating a drop of dye in the middle of our
27          # simulated region
28          block_low = int(grid_shape[0] * 0.4)
29          block_high = int(grid_shape[0] * 0.5)
30          for i in range(block_low, block_high):
31              for j in range(block_low, block_high):
32                  grid[i][j] = 0.005
33
34          # Evolve the initial conditions
35          for i in range(num_iterations):
36              grid = evolve(grid, 0.1)

```


mprof plotted



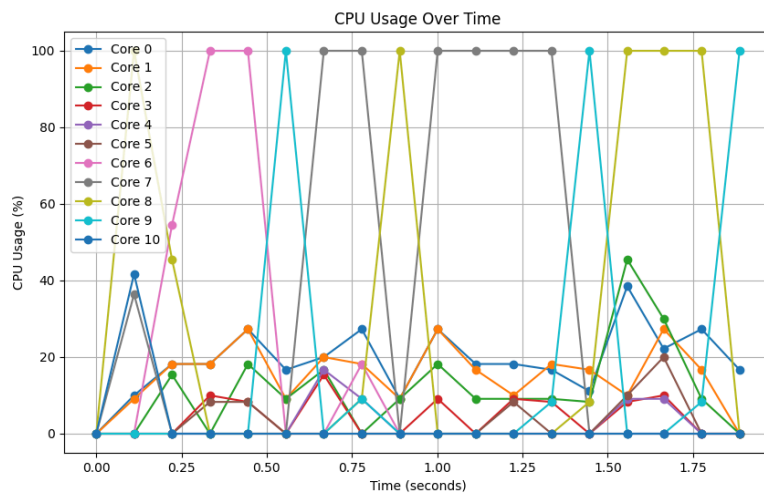
Bonus Exercise:

We made a special class in Python to make it easy to call `with` statement and the wanted interval in seconds. We wanted it to not interfere with the running code so we made a secondary thread using the module "threading" as making that thread sleep was a simple way to handle periodic check of the core usage percentage.

It takes periodic check and saves it to a list called samples. When the given `with` statement is finished, the `__exit__` function ends the thread and makes it possible for calling `profiler.plot()` and `profiler.summary_table()` to view the plot of the data and summary table. We did not include minimum values in the summary table as they always had in our tests a value of 0% on each core. The `plot()` first creates graphs for each respective core and then a combined graph with all of the core percentages over time. We decided to only include the combined graph as the computer used has 11 cores, an unnecessary amount of screenshots.


Profiler is seen in the folder A1/B/profiler.py

Result of running JuliaSet code with our own profiler tool at interval 0.1



Core	Avg Usage	Max Usage	Deviation	Time above 50%
Core 0	15.07%	36.40%	8.40%	0.00%
Core 1	10.12%	27.30%	9.19%	0.00%
Core 2	4.43%	10.00%	4.45%	0.00%
Core 3	2.44%	9.10%	3.94%	0.00%
Core 4	1.39%	16.70%	4.17%	0.00%
Core 5	4.54%	54.50%	13.63%	5.56%
Core 6	19.78%	100.00%	31.89%	16.67%
Core 7	18.24%	100.00%	32.06%	16.67%
Core 8	20.53%	100.00%	31.28%	11.11%
Core 9	41.11%	100.00%	43.10%	44.44%
Core 10	10.11%	83.30%	21.16%	5.56%

Result of running diffusion equation code with our own profiler tool at interval 0.5



Core	Avg Usage	Max Usage	Deviation	Time above 50%
Core 0	18.89%	60.00%	15.92%	1.67%
Core 1	14.32%	60.00%	13.19%	0.56%
Core 2	8.61%	50.00%	10.62%	0.00%
Core 3	4.11%	45.50%	7.26%	0.00%
Core 4	1.95%	40.00%	4.73%	0.00%
Core 5	1.43%	46.20%	5.12%	0.00%
Core 6	17.64%	100.00%	24.82%	13.33%
Core 7	28.19%	100.00%	35.33%	21.67%
Core 8	27.79%	100.00%	35.36%	19.44%
Core 9	30.83%	100.00%	36.84%	24.44%
Core 10	11.96%	100.00%	22.96%	6.11%